

```
In [2]: import pandas as pd
data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
print(data)
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1	lounge	51	882	25000	1	
1	2	pop	51	1186	32500	1	
2	3	sport	74	4658	142228	1	
3	4	lounge	51	2739	160000	1	
4	5	pop	73	3074	106880	1	
...
1533	1534	sport	51	3712	115280	1	
1534	1535	lounge	74	3835	112000	1	
1535	1536	pop	51	2223	60457	1	
1536	1537	lounge	51	2557	80750	1	
1537	1538	pop	51	1766	54276	1	

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

```
In [3]: data.head(10)
```

```
Out[3]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
5	6	pop	74	3623	70225	1	45.000702	7.682270	7900
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
8	9	sport	73	4049	76000	1	45.548000	11.549470	5600
9	10	sport	51	3653	89000	1	45.438301	10.991700	6000

```
In [4]: data.columns
```

```
Out[4]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
              'lat', 'lon', 'price'],  
              dtype='object')
```

```
In [5]: data.shape
```

```
Out[5]: (1538, 9)
```

```
In [6]: data.describe()
```

```
Out[6]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

Removing unwanted columns

```
In [7]: data1=data.drop(columns=["ID", "lat", "lon"])
```

```
In [8]: data1
```

```
Out[8]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [9]: data1=pd.get_dummies(data1)
```

```
In [10]: data1
```

```
Out[10]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [11]: data1.shape
```

```
Out[11]: (1538, 8)
```

remove the actual value from the dataframe

```
In [12]: y=data1['price']  
x=data1.drop(columns='price')
```

In [13]:

x

Out[13]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

In [14]:

```
y
```

Out[14]:

0	8900
1	8800
2	4200
3	6000
4	5700
	...
1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

In [15]: `#!pip install scikit-learn`

split the data into training set and testing set

In [16]: `from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)# 66% and 33%`

In [17]: x_test

Out[17]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0
...
291	51	701	22000	1	1	0	0
596	51	3347	85500	1	0	1	0
1489	51	366	22148	1	0	1	0
1436	51	1797	61000	1	1	0	0
575	51	366	19112	1	1	0	0

508 rows × 7 columns


```
In [18]: x_train
```

```
Out[18]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0
...
1130	51	1127	24000	1	1	0	0
1294	51	852	30000	1	1	0	0
860	51	3409	118000	1	0	1	0
1459	51	762	16700	1	1	0	0
1126	51	701	39207	1	1	0	0

1030 rows × 7 columns

```
In [19]: y_test.head(5)
```

```
Out[19]: 481    7900
76      7900
1502    9400
669     8500
1409    9700
Name: price, dtype: int64
```

```
In [20]: y_train.head(5)
```

```
Out[20]: 527    9990  
        129    9500  
        602    7590  
        331    8750  
        323    9100  
        Name: price, dtype: int64
```

LinearRegression

```
In [21]: from sklearn.linear_model import LinearRegression  
        reg=LinearRegression()#creating object of LinearRegression  
        reg.fit(x_train,y_train)#training and fitting
```

```
Out[21]: ▾ LinearRegression  
        LinearRegression()
```

```
In [22]: y_pred=reg.predict(x_test)
```

In [23]: y_pred

```
9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,
9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,
8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,
7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,
10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,
9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,
10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,
6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,
9820.78555199, 10035.83571539, 7370.77198022, 9411.45894006,
10352.85155564, 8045.21588007, 10446.80664758, 3736.20118868,
10348.63930496, 10435.96627494, 6167.80169017, 10390.11317804,
6527.69471073, 9116.4755691 , 10484.52829 , 9335.69889855,
6709.57413543, 3390.72353093, 10106.33753331, 9792.46732008,
6239.49568346, 4996.26346266, 9044.38667681, 9868.09959448,
5484.13199252, 5698.5954821 , 10086.86206874, 8115.81693479,
10392.37800936, 6835.6573351 , 6657.61744836, 5738.50576764,
8896.80120764, 9952.37340054, 10390.28377419, 9419.10788866,
9082.56591129. 10122.82465116. 10410.00504522. 10151.77663915.
```

In [24]: `from sklearn.metrics import r2_score #to know the efficiency bw the predicted price`
`r2_score(y_test,y_pred)`

Out[24]: 0.8415526986865394

In [25]: `from sklearn.metrics import mean_squared_error#calaculating mse`
`mean_squared_error(y_test,y_pred)`

Out[25]: 581887.727391353

In [26]: `import math`
`a=581887.727391353`
`print(math.sqrt(a))`

762.8156575420782

In [27]: `y_test.head(10)`

```
Out[27]: 481      7900
         76      7900
         1502     9400
         669     8500
         1409     9700
         1414     9900
         1089     9900
         1507     9950
         970     10700
         1198     8999
         Name: price, dtype: int64
```

In [28]: `y_pred`

```
10439.02404030, 10303.81930482, 8720.0083498 , 8274.3379289 ,
6889.7195761 , 10191.45963957, 4819.0674709 , 8814.11814085,
5737.62378403, 10051.06593609, 8840.87520652, 10054.31165256,
9686.269121 , 10463.56977746, 10133.15815395, 9762.80613855,
9793.03056946, 6796.69068198, 9599.3262671 , 8488.31539047,
6705.66818403, 10307.58651641, 10045.18332239, 10120.36242166,
5836.93199112, 8772.49782933, 9680.77538859, 5719.87463854,
8398.59735084, 9680.77538859, 4334.81943405, 10015.00600846,
9850.72458719, 7864.73798641, 10072.71245374, 10552.64805598,
10253.47474908, 6861.80736606, 6484.22649656, 10374.62123623,
8426.37409382, 5447.47569851, 9914.20077691, 4687.39013431,
7885.32100747, 5431.00822998, 9911.86294348, 10390.16991322,
9680.84745901, 8844.57815539, 7764.08471024, 4257.54640953,
9882.76503303, 10341.35258769, 5736.4484335 , 10179.87154436,
9501.423448 , 7997.3181334 , 5532.33458288, 9894.57834738,
10437.97459358, 6381.35845844, 9591.23555726, 9574.27908517,
10322.30715736, 9501.22785499, 9789.955758 , 9593.26549752,
6775.82788536, 7915.34831306, 10389.98590521, 10351.58343315,
7381.32686464, 9966.53983093, 10430.87188433, 10554.43156462,
10285.85574963, 10035.88086558, 9526.63034431, 7742.78157141,
.....
```

```
In [29]: results=pd.DataFrame(columns=['Price', 'Predicted'])
results['Price']=y_test
results['Predicted']=y_pred
```

```
In [30]: results
```

```
Out[30]:
```

	Price	Predicted
481	7900	5867.650338
76	7900	7133.701423
1502	9400	9866.357762
669	8500	9723.288745
1409	9700	10039.591012
...
291	10900	10032.665135
596	5699	6281.536277
1489	9500	9986.327508
1436	6990	8381.517020
575	10900	10371.142553

508 rows × 2 columns

```
In [31]: results["Difference"]=results.apply(lambda x:x.Price-x.Predicted,axis=1)
```

```
In [32]: results.head(10)
```

```
Out[32]:
```

	Price	Predicted	Difference
481	7900	5867.650338	2032.349662
76	7900	7133.701423	766.298577
1502	9400	9866.357762	-466.357762
669	8500	9723.288745	-1223.288745
1409	9700	10039.591012	-339.591012
1414	9900	9654.075826	245.924174
1089	9900	9673.145630	226.854370
1507	9950	10118.707281	-168.707281
970	10700	9903.859527	796.140473
1198	8999	9351.558284	-352.558284

```
In [33]: results["id"]=results.index
```

In [34]: results


Out[34]:

	Price	Predicted	Difference	id
481	7900	5867.650338	2032.349662	481
76	7900	7133.701423	766.298577	76
1502	9400	9866.357762	-466.357762	1502
669	8500	9723.288745	-1223.288745	669
1409	9700	10039.591012	-339.591012	1409
...
291	10900	10032.665135	867.334865	291
596	5699	6281.536277	-582.536277	596
1489	9500	9986.327508	-486.327508	1489
1436	6990	8381.517020	-1391.517020	1436
575	10900	10371.142553	528.857447	575

508 rows × 4 columns

plot the data

```
In [35]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='id',y='Price',data=results.head(50))
sns.lineplot(x='id',y='Predicted',data=results.head(50))
plt.show()
```



In []: